

University of Science and Technology Houari Boumediene

Northwind Business Intelligence Solution

ETL Pipeline & Analytical Dashboard

Project Components:

Python ETL • Star Schema DW • Interactive Dashboards
3D OLAP Analysis • Multi-Source Integration

Student Information

Name: DAOUDI Faycal Wassim
Email: faycal.wassim.daoudi@gmail.com
Matricule: 232331740505
Level: ING 3

Contents

1 Executive Summary	3
1.1 Key Achievements - CORRECTED DATA	3
1.2 Critical Corrections	3
2 System Architecture	3
2.1 ETL Pipeline	3
2.2 Technical Stack	4
3 Data Model - Star Schema	4
3.1 Dimension Tables	4
3.2 Fact Table	4
4 ETL Implementation - CORRECTED	4
4.1 Phase 1: Extraction - FIXED	4
4.2 Phase 2: Transformation	5
4.3 Phase 3: Loading	5
4.4 Phase 4: Validation	5
5 Dashboard Analysis	5
5.1 1. Executive Summary	6
5.2 2. Employee Logistics Performance	6
5.3 3. Delivery Trend Analysis	7
5.4 4. Client Delivery Treemap	8
5.5 5. 3D OLAP Cube Analysis	9
5.6 6. Territory Distribution Sunburst	10
5.7 7. Revenue Evolution	11
6 Business Insights - CORRECTED	11
6.1 Revenue Analysis	11
6.2 Delivery Performance	11
6.3 Strategic Recommendations	12
7 Technical Challenges & Solutions	12
7.1 Challenge 1: Missing 40 Orders	12
7.2 Challenge 2: \$70,137 Revenue Gap	12
7.3 Challenge 3: Delivery Status Misclassification	12
8 Performance Metrics	12
9 Lessons Learned	13
9.1 Key Takeaways	13
9.2 Best Practices	13
10 Deployment	13
10.1 Installation	13
10.2 Expected Output	14
11 Conclusion	14
A Appendix A: SQL Schema	15
B Appendix B: Validation Tests	15

C Appendix C: Dashboard Code**16****List of Figures**

1	Four-Phase ETL Architecture	3
2	Star Schema Architecture	4
3	Executive Summary Dashboard - KPIs and Delivery Metrics	6
4	Employee Delivery Performance - Stacked Bar Chart	6
5	Delivery Trends Over Time - Area Chart	7
6	Client Delivery Analysis - Hierarchical Treemap	8
7	3D OLAP Cube - Multi-dimensional Revenue Analysis	9
8	Territory Distribution - Sunburst Diagram	10
9	Revenue Evolution - Bar and Line Chart with Trend	11

List of Tables

1	Data Discrepancy Analysis	3
2	Technology Components	4
3	ETL Execution Times	12
4	Dashboard Render Times	13

1 Executive Summary

This project implements a comprehensive Business Intelligence solution for **Northwind Traders**, addressing data fragmentation across SQL Server (13 tables) and MS Access (20 tables) through an automated Python ETL pipeline, star schema data warehouse, and interactive dashboards.

1.1 Key Achievements - CORRECTED DATA

Validated Results:

- **Total Revenue:** \$1,335,930 (corrected from \$1,265,793)
- **Total Orders:** 870 (not 830)
- **Delivered:** 841 orders (96.7%)
- **Pending:** 29 orders (3.3%)
- **Clients:** 91 unique customers
- **Employees:** 9 with territory assignments
- **Timeline:** 11 years (1996-2006)

1.2 Critical Corrections

Table 1: Data Discrepancy Analysis

Metric	Initial	Corrected	Variance
Total Revenue	\$1,265,793	\$1,335,930	+\$70,137 (+5.5%)
Total Orders	830	870	+40 (+4.8%)
Delivered	817 (98.4%)	841 (96.7%)	+24
Pending	13 (1.6%)	29 (3.3%)	+16

Root Cause: Initial ETL excluded orders without shipped dates, missing 40 pending orders and \$70K revenue.

2 System Architecture

2.1 ETL Pipeline

- Phase 1: EXTRACTION → SQL Server + Access → CSV
- Phase 2: TRANSFORMATION → Clean, harmonize, calculate
- Phase 3: LOADING → Dimensions + Facts → Parquet
- Phase 4: VISUALIZATION → 7 interactive dashboards

Figure 1: Four-Phase ETL Architecture

2.2 Technical Stack

Table 2: Technology Components

Layer	Technology
Language	Python 3.14.2
Processing	Pandas, NumPy
Database	SQL Server, MS Access
Connectivity	SQLAlchemy, PyODBC
Visualization	Plotly, Matplotlib
Storage	Parquet (59.6% compression)

3 Data Model - Star Schema




Figure 2: Star Schema Architecture

3.1 Dimension Tables

DimDate: 2,223 records (daily grain, 1996-2006)

DimClient: 91 customers with geographic attributes

DimEmployee: 9 employees with territory assignments

3.2 Fact Table

FactSales: 870 order records with measures (quantity, price, discount, total) and delivery status (Livrée/Non Livrée)

4 ETL Implementation - CORRECTED

4.1 Phase 1: Extraction - FIXED

CRITICAL FIX: Remove WHERE clauses that excluded pending orders

Listing 1: Corrected Extraction

```

1 def extract_orders():
2     # WRONG: WHERE ShippedDate IS NOT NULL (excludes pending!)
3     # CORRECT: Include ALL orders
4     query = "SELECT * FROM Orders"

```

```

5     df = pd.read_sql(query, engine)
6     print(f"Extracted {len(df)} orders")    # Must be 870
7     return df

```

4.2 Phase 2: Transformation

Key Transformations:

1. Column normalization (lowercase, strip)
2. SQL Server + Access merge (deduplicate on primary keys)
3. Delivery status: NULL/empty shipped_date = "Non Livrée"
4. Revenue calculation: price × qty × (1-discount)
5. Territory enrichment via joins

Listing 2: Delivery Status Logic

```

1 def classify_delivery(shipped_date):
2     if pd.isna(shipped_date) or str(shipped_date).strip() == '':
3         return 'Non Livr e'
4     return 'Livr e'
5
6 fact['delivery_status'] = fact['shipped_date'].apply(classify_delivery)
7 # Result: 841 Livr e, 29 Non Livr e

```

4.3 Phase 3: Loading

Parquet format: 59.6% compression, 4× faster reads vs CSV

4.4 Phase 4: Validation

Listing 3: Automated Validation

```

1 def validate_warehouse():
2     fact = pd.read_parquet('FactSales.parquet')
3
4     # Test 1: Order count
5     assert len(fact) == 870, f"Expected 870, got {len(fact)}"
6
7     # Test 2: Revenue
8     total = fact['total_amount'].sum()
9     assert abs(total - 1335930) < 10
10
11    # Test 3: Delivery distribution
12    delivered = (fact['delivery_status'] == 'Livr e').sum()
13    pending = (fact['delivery_status'] == 'Non Livr e').sum()
14    assert delivered == 841 and pending == 29
15
16    print("All validations passed")

```

5 Dashboard Analysis

5.1 1. Executive Summary



Figure 3: Executive Summary Dashboard - KPIs and Delivery Metrics

KPIs: Total revenue (\$1.34M), orders (870), delivery rate (96.7%)

Features: Year/month filter, gauge visualizations, color coding

5.2 2. Employee Logistics Performance




Figure 4: Employee Delivery Performance - Stacked Bar Chart

Horizontal stacked bar chart showing delivered (blue) vs pending (red) by employee.

Top Performers:

- Margaret Peacock: 157 orders (18.0%), 98.7% delivery
- Janet Leverling: 125 orders, 100% delivery
- Nancy Davolio: 123 orders, 8 pending (needs attention)

5.3 3. Delivery Trend Analysis




Figure 5: Delivery Trends Over Time - Area Chart

Time-series area chart revealing:

- Peak: 1997-1998 (70+ orders/month in Apr 1998)
- Dormant: 1999-2005 (near-zero activity)
- Resumption: 2006 (10-15 orders/month)

5.4 4. Client Delivery Treemap

Delivery Status Distribution (All Years)




Figure 6: Client Delivery Analysis - Hierarchical Treemap

Hierarchical blocks sized by order count, colored by status. Major clients: Save-a-lot Markets, QUICK-Stop, Ernst Handel.

5.5 5. 3D OLAP Cube Analysis

3D OLAP:Year x Client x Employee




Figure 7: 3D OLAP Cube - Multi-dimensional Revenue Analysis

Innovation: Gap markers show zero-revenue (Time \times Client \times Employee) combinations.

Listing 4: Gap Analysis Implementation

```

1 import itertools
2
3 # Generate complete grid
4 grid = pd.DataFrame(
5     list(itertools.product(years, clients, employees)),
6     columns=['year', 'client', 'employee']
7 )
8
9 # Merge with actuals, fill zeros
10 df_dense = pd.merge(grid, actuals, how='left').fillna(0)
11
12 # Separate gaps from real sales
13 gaps = df_dense[df_dense['revenue'] == 0]

```

```

14  actuals = df_dense[df_dense['revenue'] > 0]
15
16  # Plot both layers (gaps as light grey dots)

```

Business Value: Reveals untapped client-employee-time opportunities.

5.6 6. Territory Distribution Sunburst

Territory Distribution by Employee




Figure 8: Territory Distribution - Sunburst Diagram

Hierarchical visualization: inner ring = employees, outer ring = territories.

Coverage: Robert King leads with 10 territories.

5.7 7. Revenue Evolution

Yearly Revenue Evolution




Figure 9: Revenue Evolution - Bar and Line Chart with Trend

Bar + line combo with average baseline.

Distribution: 1997 = 46.2% of total revenue (\$617K)

6 Business Insights - CORRECTED

6.1 Revenue Analysis

- **Total:** \$1,335,930 across 870 orders
- **Average Order Value:** \$1,535.61
- **Peak Year:** 1997 (\$617K = 46.2%)
- **Concentration:** 95% revenue in 1996-1998

6.2 Delivery Performance

- **Rate:** 96.7% (above 95% benchmark)
- **Perfect Record:** Janet Leverling (0 pending)
- **Improvement Area:** Nancy Davolio (8 pending)

6.3 Strategic Recommendations

1. Address **29 pending orders immediately** - 3.3% pending indicates process issue
2. Investigate **1999-2005 dormancy** - Understand 7-year gap
3. Replicate **Janet Leverling's process** - Study 100% delivery success
4. Balance territories - Robert King: 10 territories, Margaret Peacock: 157 orders
5. Analyze **1997 peak** - Replicate factors that drove 46% of revenue

7 Technical Challenges & Solutions

7.1 Challenge 1: Missing 40 Orders

Problem: Initial ETL extracted 830 orders, dashboard showed 870

Root Cause: WHERE ShippedDate IS NOT NULL excluded pending orders

Solution: Remove all WHERE filters in extraction, apply logic in transformation

7.2 Challenge 2: \$70,137 Revenue Gap

Problem: Revenue discrepancy between initial (\$1.27M) and actual (\$1.34M)

Root Cause: Missing 40 orders = missing revenue

Solution: Validate total revenue against dashboard after each ETL run

7.3 Challenge 3: Delivery Status Misclassification

Problem: 13 pending reported, actually 29

Root Cause: Inadequate NULL/empty string handling

Solution: Multi-check classification (isna + strip + parse)

8 Performance Metrics

Table 3: ETL Execution Times

Phase	Duration	Records/sec
SQL Extraction	2.3s	378
Access Extraction	3.1s	281
Transformation	5.2s	593
Loading	1.4s	2,204
Validation	1.1s	-
Total	13.1s	66

Table 4: Dashboard Render Times

Visualization	Time
Executive Summary	0.3s
Employee Performance	0.6s
Delivery Trends	0.9s
Client Treemap	0.7s
3D OLAP Cube	1.4s
Territory Sunburst	0.5s
Revenue Evolution	0.5s
Total	4.9s

9 Lessons Learned

9.1 Key Takeaways

1. **Validate Against UI:** Dashboard revealed 40 missing orders—ETL must match visualizations exactly
2. **No Filtering in Extraction:** Apply business logic in transformation, not source queries
3. **Comprehensive NULL Handling:** Check `isna()`, empty strings, and parsing failures
4. **Automated Testing:** Unit tests prevent regression (see validation code)
5. **Revenue Precision:** Round at each calculation step to avoid floating-point drift

9.2 Best Practices

- Extract all source data without WHERE clauses
- Log record counts at each pipeline stage
- Validate totals against known dashboard metrics
- Implement referential integrity checks
- Use Parquet for 60% compression + 4× speed

10 Deployment

10.1 Installation

Listing 5: Setup Steps

```

1 # Clone repository
2 git clone https://github.com/root-wassim/BI.git
3 cd BI/Northwind

4

5 # Create virtual environment
6 python -m venv venv
7 venv\Scripts\activate # Windows
8 source venv/bin/activate # Linux/Mac

9

10 # Install dependencies
11 pip install -r requirements.txt

12

13 # Configure database (edit scripts/config.py)
14 SERVER_NAME = r'.\SQLEXPRESS'
15 DATABASE_NAME = 'Northwind'
```

```
16
17 # Run ETL
18 cd scripts
19 python Main.py
20
21 # Launch dashboard
22 jupyter notebook ../notebooks/dashboard_analysis.ipynb
```

10.2 Expected Output

```
1 NORTHWIND ETL PIPELINE - STARTING
2   SQL Server: Orders -> 870 rows
3   Access: Orders -> 870 rows
4   Creating DimDate... 2223 rows
5   Creating FactSales... 870 orders
6
7 VALIDATION:
8   Total Revenue: $1,335,930.00
9   Delivered: 841 (96.7%)
10  Pending: 29 (3.3%)
11
12 ETL COMPLETED in 13.1 seconds
```

11 Conclusion

This corrected report documents a validated BI solution that successfully:

1. Integrated 870 orders from heterogeneous sources (SQL Server + Access)
2. Corrected \$70,137 revenue discrepancy through validation
3. Achieved 96.7% delivery rate with transparent tracking
4. Delivered 7 interactive visualizations including 3D OLAP
5. Implemented automated testing to prevent future errors

Key Achievement: Dashboard-driven validation caught systematic ETL errors, demonstrating the importance of UI-first testing in BI projects.

"In God we trust. All others must bring data." — W. Edwards Deming

A Appendix A: SQL Schema

Listing 6: Complete Data Warehouse Schema

```

1  -- Dimensions
2  CREATE TABLE DimDate (
3      sk_date INT PRIMARY KEY,
4      full_date DATE,
5      year INT, month INT, quarter INT
6  );
7
8  CREATE TABLE DimClient (
9      sk_client INT PRIMARY KEY,
10     bk_customer_id VARCHAR(10) UNIQUE,
11     company_name VARCHAR(100),
12     city VARCHAR(50), country VARCHAR(50)
13 );
14
15 CREATE TABLE DimEmployee (
16     sk_employee INT PRIMARY KEY,
17     bk_employee_id INT UNIQUE,
18     employee_name VARCHAR(100),
19     territories VARCHAR(500)
20 );
21
22 -- Fact Table
23 CREATE TABLE FactSales (
24     fact_id INT PRIMARY KEY,
25     bk_order_id INT,
26     sk_date INT REFERENCES DimDate(sk_date),
27     sk_client INT REFERENCES DimClient(sk_client),
28     sk_employee INT REFERENCES DimEmployee(sk_employee),
29     quantity INT CHECK (quantity > 0),
30     unit_price DECIMAL(10,2) CHECK (unit_price >= 0),
31     discount DECIMAL(5,2) CHECK (discount BETWEEN 0 AND 1),
32     total_amount DECIMAL(12,2),
33     delivery_status VARCHAR(20) CHECK (delivery_status IN ('Livr e', 'NonLivr e'))
34 );

```

B Appendix B: Validation Tests

Listing 7: Unit Test Suite

```

1 import unittest
2
3 class TestNorthwindDW(unittest.TestCase):
4     def test_order_count(self):
5         self.assertEqual(len(fact_sales), 870)
6
7     def test_revenue(self):
8         total = fact_sales['total_amount'].sum()
9         self.assertAlmostEqual(total, 1335930, delta=10)
10
11    def test_delivery_distribution(self):
12        delivered = (fact_sales['delivery_status']=='Livr e').sum()
13        pending = (fact_sales['delivery_status']=='NonLivr e').sum()

```

```

14     self.assertEqual(delivered, 841)
15     self.assertEqual(pending, 29)
16
17     def test_referential_integrity(self):
18         fact_clients = set(fact_sales['sk_client'])
19         dim_clients = set(dim_client['sk_client'])
20         self.assertEqual(fact_clients - dim_clients, set())

```

C Appendix C: Dashboard Code

Listing 8: Executive Summary Visualization

```

1 import plotly.graph_objects as go
2
3 def create_kpi_dashboard(df):
4     total_revenue = df['total_amount'].sum()
5     total_orders = len(df)
6     delivered = (df['delivery_status']=='Livr e').sum()
7
8     fig = go.Figure()
9
10    # Revenue indicator
11    fig.add_trace(go.Indicator(
12        mode="number",
13        value=total_revenue,
14        number={'prefix': "$", 'valueformat': ",.0f"})
15    )
16
17    # Delivery gauge
18    fig.add_trace(go.Indicator(
19        mode="gauge+number",
20        value=delivered/total_orders*100,
21        gauge={'axis': {'range': [0,100]}, 'bar': {'color': "green"}}
22    ))
23
24    return fig

```

End of Report

*Northwind Business Intelligence Solution - Corrected & Validated
Engineering School ING3 - 2024/2025*